

# rktnetes

Integrating rkt and Kubernetes

Derek Gonyeo, CoreOS



# About Me

Derek Gonyeo  
rkt Scientist at CoreOS  
Student at RIT  
@dgonyeo



## Outline

- ▶ What is rkt?
- ▶ What is kubernetes?
- ▶ What is rktnetes?
- ▶ How does rktnetes work?
- ▶ Why rktnetes?



Core OS

## Outline

- ▶ **What is rkt?**
- ▶ What is kubernetes?
- ▶ What is rktnetes?
- ▶ How does rktnetes work?
- ▶ Why rktnetes?

# What is rkt?

rkt...

# What is rkt?

rkt...

- ▶ is an open source project

# What is rkt?

rkt...

- ▶ is an open source project
- ▶ is a golang binary



Core OS

# What is rkt?

rkt...

- ▶ is an open source project
- ▶ is a golang binary
- ▶ has a simple command line interface



Core OS



# What is rkt?

rkt...

- ▶ is an open source project
- ▶ is a golang binary
- ▶ has a simple command line interface
- ▶ is packaged for CoreOS, Debian, Ubuntu, Fedora, Arch, Gentoo, and Nix



# What is rkt?

rkt is a container runtime.

This means that it will:

# What is rkt?

rkt is a container runtime.

This means that it will:

- ▶ Fetch container images

# What is rkt?

rkt is a container runtime.

This means that it will:

- ▶ Fetch container images
- ▶ Cryptographically verify container images

# What is rkt?

rkt is a container runtime.

This means that it will:

- ▶ Fetch container images
- ▶ Cryptographically verify container images
- ▶ Manage a local store of container images



Core OS

# What is rkt?

rkt is a container runtime.

This means that it will:

- ▶ Fetch container images
- ▶ Cryptographically verify container images
- ▶ Manage a local store of container images
- ▶ Run container images

# What is rkt?

rkt is a container runtime.

This means that it will:

- ▶ Fetch container images
- ▶ Cryptographically verify container images
- ▶ Manage a local store of container images
- ▶ Run container images
- ▶ Provide information about running containers

# What is rkt?

Notes on rkt's architecture:



# What is rkt?

Notes on rkt's architecture:

- ▶ Has a pod-first design

# What is rkt?

Notes on rkt's architecture:

- ▶ Has a pod-first design
- ▶ Does not have a central daemon

# What is rkt?

Notes on rkt's architecture:

- ▶ Has a pod-first design
- ▶ Does not have a central daemon
- ▶ Has swappable isolation models (more on this later)

# What is rkt?

rkt is:

- ▶ Secure
- ▶ Standards Compliant
- ▶ Composable

# What is rkt?

rkt is:

- ▶ **Secure**
- ▶ Standards Compliant
- ▶ Composable

# What is rkt?

## How is rkt secure?

# What is rkt?

How is rkt secure?

- ▶ Requires images to be GPG signed

# What is rkt?

How is rkt secure?

- ▶ Requires images to be GPG signed
- ▶ Uses selinux to prevent breakouts



Core OS



# What is rkt?

## How is rkt secure?

- ▶ Requires images to be GPG signed
- ▶ Uses selinux to prevent breakouts
- ▶ Uses seccomp to drop privileges



Core OS

# What is rkt?

## How is rkt secure?

- ▶ Requires images to be GPG signed
- ▶ Uses selinux to prevent breakouts
- ▶ Uses seccomp to drop privileges
- ▶ Uses the TPM to make a tamper-proof audit log



Core OS

# What is rkt?

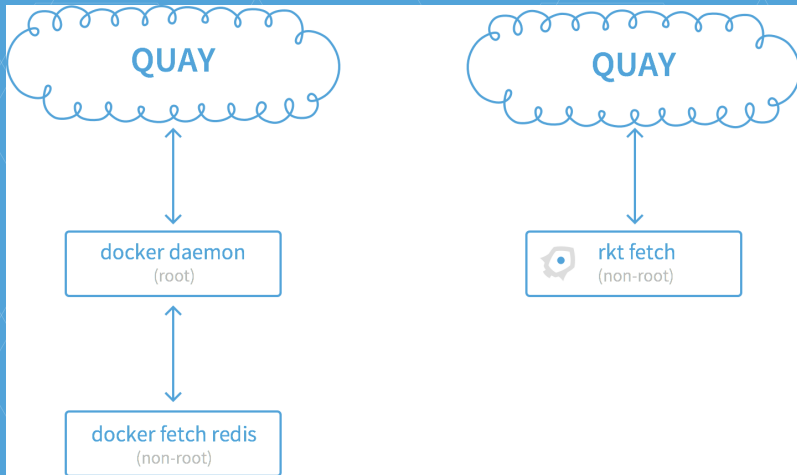
## How is rkt secure?

- ▶ Requires images to be GPG signed
- ▶ Uses selinux to prevent breakouts
- ▶ Uses seccomp to drop privileges
- ▶ Uses the TPM to make a tamper-proof audit log
- ▶ No central privileged daemon



Core OS

# What is rkt?



Core OS

# What is rkt?

rkt is:

- ▶ **Secure**
- ▶ Standards Compliant
- ▶ Composable



Core OS

# What is rkt?

rkt is:

- ▶ Secure
- ▶ **Standards Compliant**
- ▶ Composable

# What is rkt?

How is rkt standards compliant?

# What is rkt?

How is rkt standards compliant?

- ▶ Implementation of the AppC spec



# What is rkt?

How is rkt standards compliant?

- ▶ Implementation of the AppC spec
- ▶ Supports images from the docker spec

# What is rkt?

How is rkt standards compliant?

- ▶ Implementation of the AppC spec
- ▶ Supports images from the docker spec
- ▶ Supports images from the OCI spec



Core OS

# What is rkt?

How is rkt standards compliant?

- ▶ Implementation of the AppC spec
- ▶ Supports images from the docker spec
- ▶ Supports images from the OCI spec
- ▶ Native OCI support coming



Core OS

# What is rkt?

How is rkt standards compliant?

- ▶ Implementation of the AppC spec
- ▶ Supports images from the docker spec
- ▶ Supports images from the OCI spec
- ▶ Native OCI support coming
- ▶ Uses CNI for network configuration



Core OS

# What is rkt?

rkt is:

- ▶ Secure
- ▶ **Standards Compliant**
- ▶ Composable

# What is rkt?

rkt is:

- ▶ Secure
- ▶ Standards Compliant
- ▶ **Composable**

# What is rkt?

## How is rkt composable?

# What is rkt?

How is rkt composable?

- ▶ Simple process model



# What is rkt?

How is rkt composable?

- ▶ Simple process model
- ▶ Swappable stage1s

# What is rkt?


How is rkt composable?

- ▶ Simple process model
- ▶ Swappable stage1s
- ▶ Systemd integration


# What is rkt?

## Container Engine Process Models


rkt 1.0+



Docker < 1.11.0



Docker 1.11.0+



Core OS

# What is rkt?

rkt has stages

# What is rkt?

stage0:

# What is rkt?

stage0:

- ▶ Handles user interaction

# What is rkt?

stage0:

- ▶ Handles user interaction
- ▶ Fetches and verifies images

# What is rkt?

stage0:

- ▶ Handles user interaction
- ▶ Fetches and verifies images
- ▶ Manages the image store



# What is rkt?

stage0:

- ▶ Handles user interaction
- ▶ Fetches and verifies images
- ▶ Manages the image store
- ▶ Handles dependencies

# What is rkt?

stage0:

- ▶ Handles user interaction
- ▶ Fetches and verifies images
- ▶ Manages the image store
- ▶ Handles dependencies
- ▶ Renders images



Core OS

# What is rkt?

stage0:

- ▶ Handles user interaction
- ▶ Fetches and verifies images
- ▶ Manages the image store
- ▶ Handles dependencies
- ▶ Renders images
- ▶ Execs stage1



Core OS

# What is rkt?

stage1:

# What is rkt?

stage1:

- ▶ Sets up isolation for pod

# What is rkt?

stage1:

- ▶ Sets up isolation for pod
- ▶ Sets up relevant networking

# What is rkt?

stage1:

- ▶ Sets up isolation for pod
- ▶ Sets up relevant networking
- ▶ Sets up mounts



Core OS

# What is rkt?

stage1:

- ▶ Sets up isolation for pod
- ▶ Sets up relevant networking
- ▶ Sets up mounts
- ▶ Execs stage2



Core OS



# What is rkt?

The stage2 is user provided - your app!

# What is rkt?

Available stage1s:

# What is rkt?

Available stage1s:

- ▶ coreos



# What is rkt?

Available stage1s:

- ▶ coreos
- ▶ host

# What is rkt?

Available stage1s:

- ▶ coreos
- ▶ host
- ▶ fly

# What is rkt?

Available stage1s:

- ▶ coreos
- ▶ host
- ▶ fly
- ▶ kvm



Core OS

# What is rkt?

systemd isn't required to use rkt!

# What is rkt?

systemd isn't required to use rkt!

But if you are using systemd:

- ▶ Trivial to have rkt in systemd units



Core OS



# What is rkt?

systemd isn't required to use rkt!

But if you are using systemd:

- ▶ Trivial to have rkt in systemd units
- ▶ journald integration



Core OS

# What is rkt?

systemd isn't required to use rkt!

But if you are using systemd:

- ▶ Trivial to have rkt in systemd units
- ▶ journald integration
- ▶ machined integration

## Outline

- ▶ **What is rkt?**
- ▶ What is kubernetes?
- ▶ What is rktnetes?
- ▶ How does rktnetes work?
- ▶ Why rktnetes?

## Outline

- ▶ What is rkt?
- ▶ **What is kubernetes?**
- ▶ What is rktnetes?
- ▶ How does rktnetes work?
- ▶ Why rktnetes?

# What is kubernetes?


Kubernetes is:

- ▶ container orchestration
- ▶ cluster-level
- ▶ fault tolerant

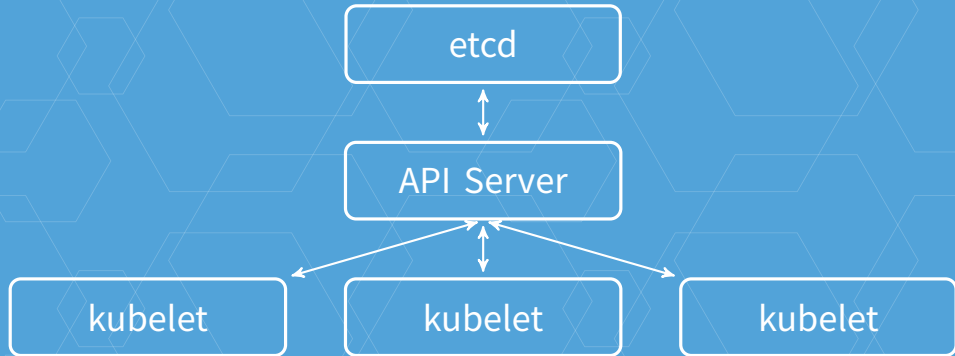
# What is kubernetes?

etcd

# What is kubernetes?



# What is kubernetes?






# What is kubernetes?

kubelet

# What is kubernetes?



# What is kubernetes?



## Outline

- ▶ What is rkt?
- ▶ **What is kubernetes?**
- ▶ What is rktnetes?
- ▶ How does rktnetes work?
- ▶ Why rktnetes?

## Outline

- ▶ What is rkt?
- ▶ What is kubernetes?
- ▶ **What is rktnetes?**
- ▶ How does rktnetes work?
- ▶ Why rktnetes?

# What is rktnetes?

Use rkt as the container runtime!

# What is rktnetes?

Current status:

# What is rktnetes?

Current status:

- ▶ Support included in Kubernetes 1.3



Core OS



# What is rktnetes?

## Current status:

- ▶ Support included in Kubernetes 1.3
- ▶ >90% of common end-to-end tests pass with rkt



Core OS

# What is rktnetes?

## Current status:

- ▶ Support included in Kubernetes 1.3
- ▶ >90% of common end-to-end tests pass with rkt
- ▶ Still working on refining the integration


## Outline

- ▶ What is rkt?
- ▶ What is kubernetes?
- ▶ **What is rktnetes?**
- ▶ How does rktnetes work?
- ▶ Why rktnetes?

## Outline

- ▶ What is rkt?
- ▶ What is kubernetes?
- ▶ What is rktnetes?
- ▶ **How does rktnetes work?**
- ▶ Why rktnetes?

# How does rktnetes work?




# How does rktnetes work?

kubelet

# How does rktnetes work?

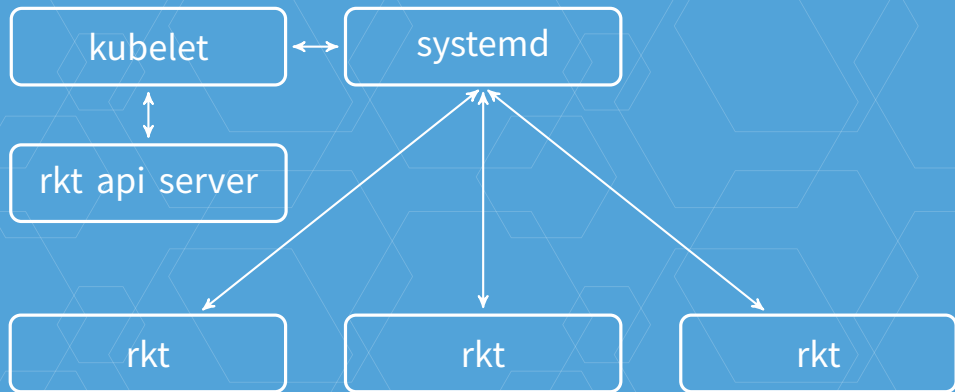


# How does rktnetes work?





# How does rktnetes work?



# How does rktnetes work?

rktnetes prerequisites:

# How does rktnetes work?

rktnetes prerequisites:

- ▶ The host must use systemd

# How does rktnetes work?

rktnetes prerequisites:

- ▶ The host must use systemd
- ▶ rkt v1.9.1 or greater must be installed



Core OS

# How does rktnetes work?

rktnetes prerequisites:

- ▶ The host must use systemd
- ▶ rkt v1.9.1 or greater must be installed
- ▶ The rkt api service must be running



Core OS

# How does rktnetes work?

## How do I use rktnetes?

# How does rktnetes work?

How do I use rktnetes?

- ▶ Add `--container-runtime=rkt` to the kubelet



Core OS

## Outline

- ▶ What is rkt?
- ▶ What is kubernetes?
- ▶ What is rktnetes?
- ▶ **How does rktnetes work?**
- ▶ Why rktnetes?



## Outline

- ▶ What is rkt?
- ▶ What is kubernetes?
- ▶ What is rktnetes?
- ▶ How does rktnetes work?
- ▶ **Why rktnetes?**

# Why rktnetes?

## Why did we build this?

- ▶ Path to a generic container runtime interface
- ▶ rkt has a different feature set than docker



Core OS

Questions?



By the way:



By the way:

# Kubernetes Community Bash

Tonight at 7 pm

Real Sports Bar & Grill

<https://kubernetesecosystembash.eventbrite.com/>



Core OS